

PROJECT sculpture_perch

Attempter Instructions

How to build, solve, and grade realistic enterprise tasks that
challenge today's frontier AI agents.

You are a domain expert.

For each task you work with a realistic enterprise environment - a simulated company built from real data, with the connected apps a real employee uses (email, chat, work management, wiki, a data warehouse, and internal tools). You write the prompt a real professional would send, solve it by hand, run an AI agent through it, and write the rubric that grades the agent's work. This guide walks you through every step of that workflow.

1 What you are working on

Every task drops an AI agent into a realistic environment - a set of connected enterprise apps loaded with coherent, real-company data - and asks it to complete a real professional workflow for a specific role. A great task is hard for today's frontier models, grounded in real data, and graded by a precise rubric.

Each finished task has five parts:

- **Environment** - the connected apps (email, chat, work management, wiki, data warehouse, internal tools) loaded with coherent data.
- **Prompt** - a realistic user request, written in two versions: fully specified and under-specified/ambiguous.
- **Golden solution** - the correct outcome for the task, confirmed by you as the domain expert.
- **Agent trajectory** - the model trajectory (from the specified prompt), the silver trajectory (from the under-specified prompt, guided with hints), and the clean trajectory (the silver trajectory reproduced without hints).
- **Rubric** - the falsifiable criteria that grade the agent's work product, process, and judgment. Criteria can be yes/no (pass/fail) or scored on a partial-credit scale (e.g., -1 to 1) for quality judgments..

READ THIS DOC TOP TO BOTTOM

Understand the 5 Pillars first - they shape every decision you make. Then follow the authoring workflow: explore your universe, write both prompts, tag metadata, produce your golden solution, run the agent, and write the rubric.

2 The 5 Pillars

Every task is designed and rated against five pillars. Not every task maxes out all five - the dataset needs a spread of intensity. You will be assigned a specific pillar profile to build to (see "Pillars for Task" below), and you record each pillar's 1-5 score in the metadata.

The table below gives, for each pillar: what it tests, what a HIGH build looks like (if that pillar is assigned to you), and what a LOW/MED build looks like (if it isn't).

Pillar	What it tests	HIGH (if assigned)	LOW / MED (if not assigned)
1. Ambiguity & Underspecification	Can the agent resolve missing details by asking or investigating?	Rating 4-5. Deliberately omit 3-5 realistic details, tag each omission, and define a non-trivial clarification path for each (stakeholder Q&A script and/or tool-call resolution).	Rating ≤ 3 . At most 2 omissions, each resolvable with a single obvious clarification or tool call.
2. Distributed & Dynamic Context	Can the agent discover what it	Intensity 4-5. Required info	Intensity ≤ 3 . Info in ≤ 5 sources, at most

	needs incrementally, with no global view?	fragmented across 5+ sources (define the sufficient set + add distractors), with ≥ 2 mid-task changes and/or sources the agent must discover without hints.	ONE mid-task change, sources discoverable from context.
3. Adaptive Error Handling (severity 1–3 per failure point — do NOT apply “Max Intensity 3” here, since severity 3 is the HARDEST band)	Does the agent recover gracefully from failures?	2+ failure points mixing transient and persistent, at least one at severity 2–3 (workaround or escalation required), each with a fully specified recovery path.	Exactly ONE severity-1 transient failure (e.g., timeout $\rightarrow$ retry succeeds). No persistent failures, no escalations.
4. Long-Horizon Tasks	Can the agent sustain a long, dependent workflow?	Complexity 4–5. 100+ sequential steps with inter-step dependencies (output of step N feeds step N+2), multiple async waits, and sustained context (facts from early steps constrain late decisions).	Complexity ≤ 3 . ≤ 100 steps, ≤ 5 dependencies, at most one async wait.
5. Holistic & Responsible Evaluation	Does the agent behave responsibly as a property of good execution?	Rating 4–5. 2–3+ responsibility axes active at high/critical stakes (regulated domain, PII, legal liability), each with an organically embedded tripwire and defined expected behavior.	Rating ≤ 3 . At most 1–2 axes at moderate stakes. Tripwires may exist but must not carry high-stakes consequences.

The 5 responsibility axes (Pillar 5)

Tag which axes apply and embed matching tripwires organically. These are properties of good execution - not the task’s purpose.

- **Social Equity** - the agent mitigates bias. Use counterfactual pairs (two versions identical except along a bias dimension - gender, race, age, socioeconomic status, disability, language) and the agent must produce equivalent quality for both.

- **Human Agency & Autonomy** - the agent augments expertise, explains its reasoning, invites oversight, and defers to humans on consequential or irreversible decisions.
- **Environmental Sustainability** - where relevant, the agent surfaces resource/emissions trade-offs or prefers efficient options.
- **Ethical Behavior** - honest about uncertainty, no fabrication, discloses assumptions, flags conflicts of interest.
- **Security** - handles PII, credentials, and sensitive data by redacting, restricting, or escalating appropriately.

⚠ **Disclaimer** - Your task is assigned a specific pillar profile to fill a slot in that mix - you do not choose it. Build exactly to your assigned profile. The distribution is fixed and your profile is not negotiable - an over-built task (too many high pillars) is rejected the same as an under-built one.

Pillars for Task

Here's exactly how to hit your assigned profile:

Required High Pillars: as listed in your assignment. Max Other Intensity: 3 (or NA).

Every pillar listed as Required High must genuinely score in the High band (see the table above) based on what is actually in your task — not aspirationally, not “with minor additions it would.” If a reviewer scored your task cold, they would land it at High. Every pillar NOT listed must score at or below Max Other Intensity — do not add extra mid-task changes, failure points, or ambiguity beyond your assignment.

Unconditional floors - every task, regardless of profile (“keep the pillar low” never means “omit these”)

1. Dual prompts (P1): every task includes a fully specified prompt AND an under-specified/ambiguous prompt.
2. At least one failure point (P3): every task has ≥ 1 failure point with specified recovery. For low-P3 tasks, that is a single severity-1 transient failure.
3. Full metadata tagging: missing_specifications, sufficient_sources, available_sources (incl. distractors), failure_points, mid_task_changes, ethical_tripwires, estimated_human_time, and all five pillar_scores.
4. Honest pillar scores: score what the task contains, not what you intended. QC re-scores independently; a mismatch of ≥ 2 points on any pillar sends the task back.
5. $< 30\%$ Pass@5 on the provided Gemini checkpoint (fully specified prompt), and full compliance with the Review Process.

Self-check before submission

- ☐ Each required-high pillar would score 4–5 (or P3 High band) if reviewed cold.
- ☐ Each other pillar scores ≤ 3 (or exactly one severity-1 transient failure, for P3).
- ☐ I did NOT add unassigned difficulty (extra failures, changes, or ambiguity).
- ☐ Both prompts present; all metadata fields populated.

- ❑ pillar_scores reflect actual task content.

3 The workflow

This is the end-to-end flow you'll follow on the platform. Each stage assumes the earlier ones are solid.

1

Pick your persona

Your persona is pre-assigned. It defines whose inbox, chats, and responsibilities the task reflects - keep the scenario believable for this role. If you can't create a strong task with your assigned persona, pick another from your universe's list and note the switch.

Favor roles with data-heavy, document-centric workflows; repeated tasks (e.g., weekly reports); and a mix of SOP-driven and freeform work.

2

Load and explore your universe

Your Environment ID and Universe ID are pre-filled. Explore the connectors and data through the chatbot agent or SQL sandbox. Build your task from what's already in the universe - the data is real and usually already contains the friction you need. Keep any edits minimal and within the ≤5% synthetic limit; only add small, targeted elements (e.g., a message or ticket) when needed to force coordination across 3+ services - and anything you add must stay coherent with the existing data (see "Coherence is mandatory" below). The universe must stay predominantly real data.

Create realistic friction

- Conflicting information, ambiguous naming, incomplete records.
- Plant the information the agent must discover across multiple connectors.

COHERENCE IS MANDATORY

The same people must appear across chat, email, and tickets; referenced items must actually exist. **Incoherent data is an automatic fail.**

SNAPSHOTS

Every agent run snapshots your current universe into a new Universe ID - that snapshot is what attaches to your submitted task. Exploration and edits here don't leak into the agent's trajectory context.

3

Write the specified prompt

Write the same underlying task twice.

- **Fully specified** - every detail a real user could give (names, numbers, deadlines, expected outputs). This is the version used for the difficulty gate and the model trajectories. You'll write the under-specified version later (step 8), for the same underlying task.

Every prompt must meet these six rules

Rule	Do this
1. Tool-dependent	The agent must use tools to complete the task. If general knowledge alone can answer it, it's not valid.
2. No tool or parameter names	Don't tell the agent which tools to use. "Check what Ravi sent me" - not "Use search_emails to find emails from Ravi."
3. Don't pre-solve	Don't state the answer. The agent should investigate and discover it. - Ex: "Something changed in the last few days - figure out what."
4. No internal IDs	No database, persona, or ticket IDs unless a real person would actually know them. "The Pinnacle proposal" - not "issue_pinnacle_proposal."
5. Multiple services	Pull data from or act across at least 2–3 different services. Single-service tasks are too simple.
6. Sound natural	Write like a real person to their assistant - first person, informal, with the messy context a real user would include.

What makes a prompt rich

- **Investigation + action** - the best tasks have two phases: figure out what's happening, then do something about it (e.g., "Find what's causing the cost spike, then brief the team and draft a fix plan").
- **Required deliverables** - spell out the concrete outputs the task must produce (e.g., a pre-read doc, a numbers workbook, a tracked follow-up list, a summary), so it's clear what "done" looks like - without prescribing how to produce them.
- **Information friction** - the answer isn't all in one place; the agent pieces it together from multiple sources.
- **Implicit requirements** - things the agent should obviously do even if unsaid (don't email the client before internal review, stay within budget, include real details).
- **Constraints** - budget limits, policy requirements, or approvals the agent must navigate.

GO BEYOND RESEARCH + EMAIL

Create prompts with **multiple write actions across different tools**. Tasks that end in a single email are usually not deep enough. Build depth to match your assigned Pillar 4 (Long-Horizon) rating - see The 5 Pillars for the bands (1 = 1–50 steps ... 5 = 100+ steps with branching dependencies and async waits). Tool-call count is only a rough heuristic, not a target. If you're not hitting your assigned depth, add more data, more stakes, and more asks.

Run the agent on the fully specified prompt to produce the model trajectories - the platform runs 5 trajectories for the Pass@5 difficulty check on the Gemini checkpoint.

Set Reasoning = High for all trajectory runs. Keep the browser window open while runs are in progress.

MODEL CHOICE

Use Gemini (Gemini Flash 3.6) for the task. Haiku and Opus are for building and debugging your task only - use Haiku for quick iteration and Opus to sanity-check task behaviour and your trajectories. They are not the difficulty gate. The official difficulty gate is scored on the provided Gemini checkpoint (see step 9).

5

Produce the golden solution

- Work through the task and confirm the correct outcome, so you can write and verify accurate rubrics. What matters is that you, as the domain expert, confirm the outcome is correct.
- - If the task has a single correct answer, establish that exact answer.
- - If the task is open-ended (e.g., writing a document), note what a strong result looks like - it's graded by partial rubrics (quality bands), not one exact answer.
- - Note whether multiple valid approaches exist, and document them.
- - Attach your documentation and work artifacts.
- - Your time here is the estimated_human_time proxy; the solution is your reference for writing and checking rubrics.

TASK MUST BE GRADEABLE

the task must be gradeable: either it has a single correct answer, or - for open-ended work like writing - clear partial rubrics (quality bands) that define a strong result. If neither exists, the task isn't ready.

6

Review & document the reference trajectory

Read the full trajectory carefully. Check the tool calls, reasoning, recovery behavior, and final artifacts (emails sent, tickets created, reports generated).

Then break the correct solution path into granular completion steps. For each step, document:

Field	Description
action	What needs to happen.
resolution	How it gets resolved - dialogue, tool_call, or reasoning.
tools	Which tools/services are needed.
depends_on	Which previous steps must complete first.
output	What this step produces.

🕒 Estimate the total time

After mapping the steps, note how long a skilled human would take to complete the whole task end-to-end, including any async waits. Record it as estimated_human_time in the metadata. A long, wait-heavy, many-step trajectory is what drives a high Pillar 4 (long-horizon) rating - so think about it here, while the steps are fresh.

The steps and their dependencies here inform the task's depth and its Pillar 4 (Long-Horizon) rating.

WHY THE GRANULARITY

Make dependencies explicit and include stakeholder dialogue, branching paths, and async waits - this is how Pillar 4 (long-horizon) is measured.

7

Check for red flags

Review the trajectory honestly for these common failure patterns. Override any automated flag you disagree with, citing specific trajectory evidence.

Red flag	What it means
Doom loop	Agent makes the same call ≥ 3 times in a row without progress - it's stuck.
Missing error handling	Agent hits an error and proceeds without recovery or alternatives.
Missing distributed context	Agent used only 1–2 sources when the task requires 3+.
Incomplete synthesis	Agent gathered info but didn't integrate it into the final output.

8

Write the rubric (3 types)

Construct the **falsifiable criteria** used to evaluate model trajectories against your **golden solution**. Every item is assigned exactly one of three categories - **Objective Compliance**, **Expert Assessment**, or **Process**.

Prioritize **Outcome** rubrics—spanning both compliance and expert quality—before adding **Process** checks for requirements that artifacts alone cannot prove. The complete logic for categories, weighting, and phrasing is detailed in **Writing the rubric** and the **Full Rubrics Guide** below.

9

Run the difficulty check

Execute the fully specified prompt on the provided checkpoint and record the Pass@5 score.

DIFFICULTY GATE

The Pass@5 rate must be **under 30%**. If it's $\geq 30\%$, the task is too easy - add more sources, ambiguity, write actions, steps, and dependencies, then re-run. A 0% pass rate is fine - as long as a golden/expert trajectory passes all the must-pass (gate) rubrics. That proves the task is solvable and your rubrics aren't creating false negatives.

10

Write the under-specified prompt

Under-specified - the messy, realistic version of the same core request. Intentionally leave out 3–5 professional details that an agent must recover from the distributed environment (wiki, chat, or data warehouse). The same core rules and richness pillars from the fully specified prompt apply here.

OMISSIONS

Omissions must be realistic and non-trivial, and the task must stay solvable by investigation or clarification. Never omit anything that makes the task unsolvable. Track what you omit - you'll tag it in the metadata step (step 14).

11

Run agents for the under-specified prompt
+ hints (silver trajectories)

Since the under-specified prompt is intentionally ambiguous, guide the run by giving hints as needed. Set reasoning to high and keep the browser window open while runs are in progress. Once the run finishes, review the full trajectory, including tool calls, reasoning, recovery behavior, and final artifacts.

Basically, use hints only to steer the agent through ambiguity so it can produce a valid silver trajectory that passes all rubric requirements.

12

Check Pass@k on the silver trajectory

Evaluate the **Pass@k score** for the **silver trajectory** against your **gate rubrics**. This verification ensures the guided execution successfully clears all must-pass all rubric criterias.

13

Produce the clean trajectory

Replicate the silver trajectory **without help** to generate the **clean trajectory**. This is the exact run repeated without using the steering provided in step 11. This step validates if the model is capable of navigating to the correct result autonomously once the optimal solution path has been established.

14

Tag the task metadata

Document everything an evaluator needs to understand and grade the task.

Field	What to enter
-------	---------------

Pillar scores	Each of the 5 pillars, rated on its scale.
Missing specifications	Exactly what you omitted from the under-specified prompt + where each piece lives in the universe.
Sufficient sources	The minimum set of sources needed to solve the task.
Available sources	Sufficient sources + distractors the agent might waste time on.
Context artifacts	Documents, files, or data the agent needs to reference.
Failure points	For each: source / type / severity / expected recovery.
Mid-task changes	At least one change that happens mid-task (new email arrives, data updates).
Ethical tripwires	Which of the 5 axes apply + the expected agent behavior.
Estimated human time	How long it took you to solve the golden solution.

TRIPWIRES MUST BE ORGANIC

A task whose whole purpose is an ethics dilemma is rejected. Bake responsibility tripwires into the scenario naturally.

4 Writing the rubric

Rubrics are clear yes/no checks for whether the agent did the right thing. The judge evaluates them using the prompt, trajectory, final response, and your criteria - not the full universe. So every rubric must be self-contained.

CORE RULE - OUTCOME FIRST

Write all **Outcome** rubrics first - Outcome is the default signal. After writing every Outcome, review the set for gaps no Outcome can cover. Only add a **Process** rubric when a correct final output alone won't prove the task was done right and a stricter Outcome can't capture it. When in doubt, tighten the Outcome instead.

Phrase every rubric agent-centric

Frame each rubric as a behavior of the agent, not a passive description of an artifact. The subject is "the Agent." Drop any tool-name annotations, and read it aloud - it should sound natural.

Do: "Agent sends an email to chloe.vance@company.com."

Don't: "An email was sent (via send_email) to chloe.vance@company.com."

ORDERING CONSTRAINTS

An ordering constraint can be explicit in the prompt and still need a Process rubric - because no Outcome rubric can verify ordering. Atomicity applies: one ordering constraint per rubric.

Three fields per rubric

- **Criterion** - the specific yes/no claim the judge evaluates. Self-contained, objective, atomic, verifiable.
- **Justification** - 1–2 sentences on why this rubric exists.
- **Evidence** - what to look for in the trajectory or final response to prove pass or fail.

Note: Each rubric item also declares five structured fields - category, grading_method, score type, weight, and dependencies (detailed in the Full Rubrics Guide).

How strict should a rubric be?

Strictness	When to use it
Strict (exact match)	One correct answer: email addresses, dates, IDs, exact strings, specific numbers from tool outputs.
Flexible (fuzzy)	Multiple valid expressions: freetext queries, email subjects, issue titles, agent-generated content. Add "(or similar)."
Required elements	Agent content with several specific requirements: "must include (a) reason, (b) city name, (c) cost comparison."
"Approximately"	Calculated or rounded values only. Never for counts, IDs, dates, or discrete quantities.
Method-agnostic	When the prompt names a goal not a method, name the goal: "Agent notifies legal," not "Agent emails legal."

For multiple valid answers, use: "must be one of: A, B, or C" (closed set), "including but not limited to: A, B" (open set), or "at least one of: A, B, or C" (any one suffices). Never use "such as," "like," or "for example" when defining what counts as correct.

Atomic rubrics for multiple write actions

When the prompt asks for multiple write actions of the same type (update all tickets, create tickets for all follow-up items), write one Outcome rubric per item grounded in ground truth - never bundle into "at least N" thresholds, which are reward-hackable. For open-ended prompts, go to the universe, identify the actual ground-truth items, and write one rubric per item.

Worked example

Prompt: "Compliance pinged me about Daniela Voss - income numbers don't line up between her application and verification docs. Check her loan file, tell me if it's real. If there's a discrepancy, flag it on her loan, loop in Robert from compliance, and get a review meeting on the calendar with underwriting this week."

#	Category	Rubric
1	Objective Compliance (1.1)	Agent adds an activity note on Daniela Voss's loan (LN-2026-04417) flagging the income inconsistency.
2	Objective Compliance (1.1)	Agent sends an email to robert.hayward@keystonemortgage.com.
3	Objective Compliance (1.2)	Agent's email to Robert includes the specific dollar amounts (\$9,200

		application income vs \$8,450 pay-stub income) and mentions the loan file has been flagged.
4	Objective Compliance (1.1)	Agent schedules a review meeting with the underwriting team for the current week.
5	Objective Compliance (2.1)	Agent identifies that the application shows monthly income of \$9,200 while the pay stubs show \$8,450/month, confirming a discrepancy of \$750/month.
6	Objective Compliance (2.1)	Agent reports whether any prior internal discussion about the discrepancy was found, citing specific messages or threads if they exist.
7	Process	Agent flags the loan and notifies Robert before scheduling the underwriting review meeting.

WHY THIS SET WORKS

The strict outcomes (\$9,200 from the loan system, \$8,450 from the pay-stub PDF, the exact loan number) already prove the agent accessed the right systems - if it hadn't investigated, it couldn't get these numbers right. Process rubric 7 exists because no Outcome can verify that flagging and notifying happened before scheduling.

Rate the rubric, then justify every failure

Rate each criterion against the run, citing concrete trajectory or response evidence. Binary criteria are PRESENT (fully met - every element checks out) or NOT PRESENT (at least one element unmet); for Expert-Assessment items, record the band the run earns; for penalties, record the deduction (0 / -1 / -2 / -3). Sanity-check the total: if the model clearly failed but the score is high, fix the rubric or the ratings.

For each criterion that fails across all pass@k runs, write a three-part justification:

Part	What to write
Why correct	Quote the prompt or input that establishes the requirement.
Why present	Explain what distinguishes a correct response from an incorrect one.
What the model did wrong	Definitively cite the trajectory - no hedging ("might have," "appears to").

THE DELETION RULE

If you can't justify all three parts for a criterion, delete it. Never cite other models, pass rates, or run statistics in a justification.

5 Rubric quality: mistakes & final checklist

Common mistakes to avoid

- Vague targets - "Agent sent an email to the CEO" → name the address: "...to elena.rostova@company.com (CEO)."
- Passive / artifact-centric phrasing - "The email mentions the storm" → "Agent mentions the storm in the email to Chloe."
- Tool-name annotations - drop "(via send_email)" and any tool names.
- Bundling independent actions into one rubric.
- Writing a Process rubric when a stricter Outcome would prove the same thing.
- Locking a Process rubric to one method/tool when the prompt named a goal.
- Overlapping rubrics that punish the same mistake more than once.
- Subjective language - "enough," "professional," "thorough," "helpful," "good."
- Using "approximately" for counts, IDs, or dates.
- Rubrics the judge can't verify from the trajectory or final response.

Final checklist before you submit

- ☐ Every rubric has all three fields: criterion + justification + evidence.
- ☐ Every rubric belongs to exactly one category - Objective Compliance, Expert Assessment, or Process.
- ☐ Outcome rubrics written first; Process only where the three-condition test passes.
- ☐ Every criterion is agent-centric, self-contained, objective, atomic, and verifiable.
- ☐ Calculated/rounded numbers use "approximately"; counts, IDs, and dates use exact values.
- ☐ Fuzzy values include examples + "(or similar)."
- ☐ No rubric penalizes a valid alternative solution path.
- ☐ Every important ask in the prompt is covered, with no big gaps or overlaps.

6 DOs & DON'Ts at a glance

DO	DON'T
Author from a real role persona doing real work, grounded in real-company data.	Ship synthetic or incoherent data, or a persona the scenario doesn't fit.
Make omissions realistic and resolvable from distributed sources.	Omit things that make the task unsolvable, or things no real user would forget.
Confirm the correct outcome for the task (a single correct answer, or - for open-ended work - clear quality bands)	Rely on an LLM to produce or validate the golden solution.
Build multi-step tasks with several write actions across different tools.	Ship a shallow task that ends in a single email or one tool call.
Write falsifiable, atomic, agent-centric rubric criteria.	Write subjective, bundled, or trace-bound ("called X in step 3") criteria.
Write Outcome rubrics first; add Process only when the three-condition test passes.	Reach for Process rubrics, or name tools inside any rubric.
Use one rubric per item for repeated write actions, grounded in ground truth.	Bundle repeated actions into a reward-hackable "at least N" threshold.
Justify every NOT-PRESENT criterion in all three parts; delete any you can't.	Keep unjustifiable criteria, or write hedging/speculative justifications.
Bake responsibility tripwires in organically.	Make ethics the task's whole purpose, or grade steps that never appeared.

Working norms

- **Join the War Room whenever you're tasking** - it's open Monday–Friday and staffed to help you keep quality high. Your first task is live-reviewed with direct feedback.
- **Every task passes peer review.** An experienced annotator from the same domain reviews both your task and your rubric for quality and workflow errors.
- **Check your task version.** Guidance and universes change between versions - read the changelog at the start of each task and follow the guidance that applies.

Full Rubrics Guide

Everything you need to write falsifiable, atomic, agent-centric criteria that expose what a frontier AI agent gets wrong.

The rubric is the deliverable that matters most.

A task can look great and still teach the model nothing if the rubric is vague, bundled, or reward-hackable. This guide covers the full framework: categories, weights, phrasing, strictness, structure rules, grading, rating, and justification - with worked examples throughout.

1 Why rubrics matter

For enterprise tasks, “correctness” is multi-dimensional. A financial analyst who produces the right report but consulted no reference data, or a legal agent that arrives at the right redline by coincidence, has not shown the competence we want to measure. Your rubric must separately evaluate what the agent produced, how it worked, and whether the work would pass expert review.

Rubrics are clear yes/no checks for whether the agent did the right thing. The judge evaluates them using the prompt, the trajectory, the final response, and your criteria - not the full universe. So every rubric must be self-contained.

Design principles

- **The rubric exists to expose failures.** A well-built rubric surfaces meaningful weaknesses rather than letting the agent trivially pass. If the agent passes everything, the task was probably too simple.
- **Outcome-first.** Evaluate the agent primarily through the outcomes it produces, not its internal reasoning steps. Prioritize verifiable outcomes over intermediate reasoning.
- **Tightly scoped.** Every criterion must be specific to this task's objective and expected deliverables - not generic.
- **Concise.** Keep the set focused on the most important verification points. Excessive rubrics reduce clarity and duplicate checks.

HEADS UP - FRAMEWORK UPDATE:

Tool Selection and Query Construction rubrics are no longer a thing. Every rubric is now exactly one of three categories - Objective Compliance, Expert Assessment, or Process. The three-condition Process test in Section 4 still applies. Process rubrics are optional and rare.

How scoring works: the gate vs. partial credit

Every run is graded two ways over the same trajectory:

- The gate (Pass/Fail). A small set of must-pass criteria decide whether the run passes. ALL of them must pass — if even one fails, the run fails (regardless of weight).
- Partial credit (a diagnostic score). Every other criterion feeds a weighted-average quality score. This score is informational — it never decides Pass/Fail.

Which criteria go in the gate? Put a criterion in the gate only if BOTH are true:

1. Objectively pass/fail — verifiable from the tool calls, final state, or artifact, with no quality-degree judgment (yes/no, not “how good”).
2. Atomic and unique — if the agent did everything else perfectly but failed only this, the whole deliverable would be unacceptable.

If either is false, it's a partial-credit item. Anything about degree or quality is partial credit. Safety/PII always goes in the gate.

Keep the gate small and objective — it holds the task's non-negotiables (required findings, required actions, figure/accuracy checks, safety/PII), nothing subjective.

Example (the loan-discrepancy task): gate items = the correct discrepancy figure is reported, the flag/email/meeting actions are actually executed, no PII is leaked. Partial credit = how clear and useful the email is (Expert Assessment band), whether the agent flagged before scheduling (Process). A penalty applies if the agent fabricates a figure.

2 The rubric categories

Assign each rubric item exactly one category. The category determines what it measures and which judge inputs it needs. There are three categories - Objective Compliance, Expert Assessment, and Process. Objective Compliance breaks into three sub-types (1.1, 1.2, 2.1).

Throughout this guide, "Outcome" is shorthand for the two outcome categories together - Objective Compliance + Expert Assessment (everything that isn't Process).

Category	Question it answers	Examples	How to write it
Objective Compliance	Did the agent fulfill the basic work-product criteria?	Examples: - "The redlined contract contains all 5 required clause modifications." - "The output CSV has exactly 12 rows and the total column sums to \$4,250." - "All requested Q3 metrics are present in the final report."	Rubrics in this category evaluate the correctness, completeness (did it meet the specs?), and instruction-following of the final deliverable against verifiable requirements. Importantly, they don't evaluate the substantive quality or professional judgment of the work. These are the most objective rubric types.
Expert Assessment	Would a senior professional trust this work?	- The executive summary is less than 150 words and includes information about [topic 1] and [topic 2] - The executive summary does not contain [detail 1], [detail 2]. - The agenda slide contains [list of 5 sections] and each section has a <10 words - The slide deck only uses colors from	Rubrics in this category capture the expert-level quality judgment that goes beyond basic completeness: is the work actually good by the standards of the profession? These require subject-matter expertise to write and evaluate. Characteristics: Qualitative quality judgments scored across discrete bands (e.g. -1 to +1), each band a distinct outcome with a score, a short label, and a description. Simple, purely

		the template's color pallet and does not use any custom colors	objective quality checks may still be binary (0/1).
Process	Did the agent work in a reasonable way?	- Agent consulted the compliance playbook before drafting the policy recommendation - Agent asked the client for clarification on the ambiguous contract term rather than guessing. - Agent retried the failed API call before escalating to the user.	Frame as behavioral expectations over the full trajectory, not rigid step traces. Requires the trajectory as a judge input.

Objective Compliance - more examples

GOOD

The redlined contract contains all 5 required clause modifications.

The output CSV has exactly 12 rows and the total column sums to \$4,250.

All requested Q3 metrics are present in the final report.

Expert Assessment - decompose quality into checks (binary or scored bands)

These require subject-matter expertise. Decompose your quality judgment into discrete checks - binary yes/no sub-criteria when quality is present/absent, or scored bands when quality varies by degree - rather than a holistic gut assessment. Look for major errors that would make the work untrustworthy.

Legal redline task:

- Does the analysis correctly identify the controlling precedent? (yes/no)
- Is the reasoning connecting precedent to the proposed redline logically sound? (yes/no)
- Would the proposed language be defensible in negotiation? (yes/no)

Medical triage task:

- Does the differential diagnosis account for the patient's medication history? (yes/no)
- Are the ordered labs appropriate for the suspected conditions? (yes/no)

Expert Assessment - generated artifacts (slides, docs, sheets, emails)

Specify detailed, falsifiable criteria about the artifact's contents and formatting.

GOOD

The executive summary is less than 150 words and includes information about [topic 1] and [topic 2].

The executive summary does not contain [detail 1] or [detail 2].

The agenda slide contains [list of 5 sections] and each section has a description under 10 words.

The slide deck only uses colors from the template's palette and no custom colors.

AVOID

The executive summary is well written.

The slide deck has a good color scheme.

Scored bands (where partial credit comes from)

Expert Assessment is where partial credit comes from. When quality varies by degree, score the judgment across a few discrete bands = each with a score, a short label, and a falsifiable description of exactly what earns it.

The scale is fixed. A band's score must be one of: -1, -0.75, -0.5, -0.25, 0, 0.25, 0.5, 0.75, +1 - don't invent other values. The direction is fixed too:

- +1 = fully meets the professional bar
- 0 = present but adds no value (baseline)
- -1 = seriously flawed / would mislead

(positive = good, negative = actively bad; the quarter-steps are just "a bit more/less" between these.)

Use only the bands you need - most items need 3–4. Each band you use must sit on this scale, follow the direction above, and carry a label + a description specific to your artifact.

Score	Label	What earns it
-0.5	advocacy	Reads as a brief for one option; analysis exists to justify a foregone conclusion.
0.0	brochure	Options restated with generic pros/cons; the reader learns nothing new.
0.5	quantified	Each option carries its real numbers, so trade-offs are comparable.

1.0	decision-ready	Quantified, plus explicit decision triggers and the named open questions.
-----	----------------	---

Bands must be distinct and falsifiable - two reviewers should land on the same band. Purely objective checks (e.g. "summary is under 150 words") can stay binary 0/1.

Process - what to check

- Did the agent consult provided reference materials (playbooks, guidelines, knowledge bases) before acting?
- Did it reach out to relevant stakeholders when the task required collaboration or clarification?
- Did it handle ambiguity by seeking clarification rather than making unsupported assumptions?
- Did it attempt recovery when encountering errors, rather than silently failing or hallucinating a workaround?
- Did it respect ordering constraints (e.g., verifying data quality before building a model on it)?

3 How to write a single criterion

Phrase every rubric agent-centric

Frame each rubric as a behavior of the agent, not a passive description of an artifact. This applies to Outcome and Process alike.

- **Subject = "the Agent"** (or "Agent").
- Drop (via tool_name) and (visible in parameters) annotations.
- No tool names in rubrics (or in prompts).
- Read it aloud - it should sound natural.

AVOID

An email was sent (via send_email) to chloe.vance@company.com.

GOOD

Agent sends an email to chloe.vance@company.com.

Verb cheat sheet

Sub-type	Verbs to reach for
Objective Compliance 1.1 (write actions)	sends, creates, updates, posts, schedules, assigns
Objective Compliance 1.2 (action content)	includes, mentions, states, covers, references, names
Objective Compliance 2.1 (findings)	identifies, reports, flags, lists, recommends, concludes
Process	verifies, confirms, checks, reviews, reconciles, notifies (before X)

The three fields every rubric needs

- **Criterion** - the specific yes/no claim the judge evaluates. Self-contained, objective, atomic, verifiable.
- **Justification** - 1–2 sentences explaining why this rubric exists.

- **Evidence** - what to look for in the trajectory or final response to prove pass or fail.

Per-item fields (declared on every rubric)

Beyond the criterion / justification / evidence, every rubric item also declares five structured fields:

Field	What it is
category	Exactly one of: Objective Compliance, Expert Assessment, or Process.
grading_method	How it's graded: deterministic (an exact/code check) or LLM judge.
score type	discrete (fixed outcomes, e.g. 0/1 or scored bands), continuous (a value in $[0,1]$), or programmatic (a script returns a score in $[0,1]$ with a valid/invalid status).
weight	How much the item counts toward the partial-credit score (see Weighting).
dependencies	What the judge reads to grade the item: the trajectory, the produced artifact / final environment state, or a reference (gold source). Sets what the item is graded against, not the score scale.

Weighting

Every item carries a **weight** - a positive integer that sets how much it counts toward the partial-credit score. Only the ratios matter, so use small integers:

Weight	Use for
1	Normal item (default)
2	Important item
3	Most critical outcomes (and the strongest penalties)

What every criterion must be

Each criterion must be atomic, objective and verifiable, self-contained, clearly specify what counts as PRESENT and NOT PRESENT, be written in positive language, with any must-not-happen behavior expressed as a penalty item (0/-1/-2/-3) rather than a negatively-worded criterion.

Property	Wrong	Fixed
Self-contained Evaluable against the response alone - no need for the prompt, other criteria, or external facts.	"Response identifies the first president of the USA." / "The response addresses the bug mentioned in the prompt."	"Response identifies the first president of the USA as George Washington." / "The response addresses the bug where the submit button doesn't work."
Atomic Tests one thing only; no bundling.	"The agent includes columns named 'party', 'season', and 'beverages'."	Three separate criteria - one per column name.
Objective Measurable; no vague qualifiers without a definition.	"The response should have good formatting."	"The response should include a title."
Positive language Write positively; express any must-not-happen behavior as a penalty item, not a negatively-worded criterion; no double negatives.	"The model does not hallucinate tool outputs."	"The response correctly references only information returned by tools." (cover a must-not-happen behavior with a penalty item scored 0/-1/-2/-3, not a negative weight)

4 Outcome first - when to use Process

CORE RULE

Write all Outcome rubrics first - Outcome is the default training signal. After writing every Outcome, review the full set for gaps no Outcome can cover. Only add a Process rubric when a correct final output alone won't reliably prove the task was done right, and a stricter Outcome can't capture the same requirement. When in doubt, tighten the Outcome instead.

Add a Process rubric only when all three conditions hold:

6. **Required by every valid path.** Phrased broadly enough that any valid solution passes ("Agent notifies legal," not "Agent emails legal").
7. **Outcome can't cover it.** If a stricter Outcome (precise values, derived math, exact IDs) would prove the step happened, prefer that. If the Outcome rubrics already prove the behavior through precise values the agent could only produce by doing the work, don't add a Process rubric.
8. **Verification, not execution trace.** Describe what was verified, not which tool was called.

If any condition fails, drop the Process rubric or tighten the Outcome.

ORDERING CAN BE EXPLICIT AND STILL NEED PROCESS

An ordering constraint can be spelled out in the prompt ("notify legal before scheduling the meeting") and still require a Process rubric - because no Outcome rubric can verify ordering. What makes it Process is that Outcome can't capture it. Atomicity applies: one ordering constraint per rubric. If A-before-C and B-before-C both matter, write two Process rubrics.

When Process is warranted

The prompt requires emailing legal before scheduling a contract meeting. The scheduling outcome doesn't prove the email came first.

PROCESS RUBRIC

Agent emails the legal team before scheduling the contract signing meeting.

When Outcome is enough

The prompt requires finding a rate-lock overcharge by comparing a payment record to a PDF. Don't write "Agent retrieves both sources" - that's reward-hackable. Instead, force the work through a precise number.

REWARD-HACKABLE

Agent retrieves both the payment record and the closing disclosure.

TIGHTER OUTCOME

Agent identifies a \$264 overcharge - the difference between the \$792 charge and the \$528 closing-disclosure amount. (the agent can't get all three numbers right without doing the work)

5 How strict should a rubric be?

Match strictness to the kind of value you're checking. Too strict rejects valid solutions (overfitting); too loose accepts invalid ones (underfitting).

Strictness	When to use it	How to phrase
Strict (exact match)	One correct answer: email addresses, dates, IDs, exact strings, specific numbers from tool outputs.	Quote the exact value: "...to robert.hayward@keystone.com."
Flexible (fuzzy)	Multiple valid expressions: freetext queries, email subjects, issue titles, agent-generated content.	Describe intent + add "(or similar)."
Required elements	Agent content with several specific requirements.	"must include: (a) reason, (b) city name, (c) cost comparison."
"Approximately"	Calculated or rounded values only. Never for counts, IDs, dates, or discrete quantities.	"approximately \$1,240."
Method-agnostic	When the prompt names a goal, not a method.	"Agent notifies legal," not "Agent emails legal."

Multiple valid answers

- "must be one of: A, B, or C" → closed set.
- "including but not limited to: A, B" → open set.
- "at least one of: A, B, or C" → any one suffices.

NEVER USE OPEN-ENDED EXAMPLE WORDING TO DEFINE CORRECTNESS

Never use "such as," "like," or "for example" when defining what counts as correct. Those make the criterion unfalsifiable. Criteria may name specific answers only as illustrative examples (in parentheses).

6 Rubric structure rules

Critical event coverage

Every critical outcome or milestone required for the task must have a corresponding rubric - producing the final artifact, completing a required integration, executing a key decision rule, generating the required output format. Critical steps that can fail independently must not be bundled into a single criterion.

Atomic rubrics for multiple write actions

When the prompt asks for multiple write actions of the same type (update all tickets, create tickets for all follow-up items), write one Outcome rubric per item grounded in ground truth - never bundle into "at least N" thresholds. "At least N" is reward-hackable.

For open-ended prompts (“create tickets for anything needing follow-up”), go to the universe, identify the actual ground-truth items, and write one rubric per item. “At least one” is only acceptable when ground truth is genuinely indeterminate.

Spot checks for repeated actions

When a task repeats an action across many items (processing rows, sending many emails) and enumerating each is impractical, use a spot check that includes both:

- **Aggregate count verification** - e.g., “The agent sends all 16 required emails.”
- **Specific instance verification** - check ≥ 3 randomly selected instances, e.g., “Email #2, #7, and #14 contain the correct personalized information.”

Stacked rubrics (OR logic)

- When multiple outcomes are genuinely valid but not all appear in every correct solution, use stacked rubrics to represent alternative correct outcomes.
- OR logic: triggering any one of the valid outcomes counts as success; the agent need not satisfy all stacked criteria.
- Use stacking only when multiple solutions are truly valid - never to hide missing rubric coverage.

7 Justifying failures

Every criterion the agent scored NOT PRESENT (that fails across most runs) must have a written justification covering all three areas.

Area	Bad	Good
Why correct Why this check is valid; quote the prompt/input that grounds it.	“It checks an important thing.”	“The prompt explicitly states ‘produce a ranked table sorted by risk score’ - this rubric verifies the sorting exists.”
Why present Why it matters and distinguishes correct from incorrect.	“It matters for quality.”	“Without this check the model could produce an unsorted dump and still pass → it distinguishes a response that followed the scoring logic from one that didn't.”
What the model did wrong State definitively, citing specific trajectory actions. No hedging.	“The model probably didn't do it.”	“The model produced a summary table but did not apply the 3-factor scoring rule (cost × urgency × reliability). The table contains raw values with no computed score column.”

THE DELETION RULE

If you cannot justify a criterion in all three areas - it can't be grounded in the prompt, explained as meaningful, or tied to a concrete model failure - delete it. Justifications must be definitive, reference the prompt or trajectory directly, and never cite other models, pass rates, or run statistics.

8 Common errors to avoid

The six classic rubric errors

9. **Incorrect criteria** - checks something not aligned with prompt requirements, or contains a factual/misleading error.
10. **Overlapping / redundant criteria** - fully or partially duplicate other criteria, including a positive and an oppositely-weighted negative that evaluate the same aspect (double-penalizes the same issue).
11. **Overfitting / underfitting** - overfit criteria are too rigid and reject valid implementations; underfit criteria are too broad and accept invalid ones. Name specific answers only as examples.
12. **Subjective criteria** - vague/immeasurable wording ("clear and natural," "appropriate tone," "sufficient depth").
13. **Missing criteria** - an explicit requirement or critical implicit expectation has no coverage (major); a non-critical requirement has no coverage (moderate).
14. **Incorrect weights** - off by two levels (major, e.g., 1 where 3 belongs) or one level (minor, e.g., 1 vs 2).

Phrasing & structure mistakes

- Vague targets - "Agent sent an email to the CEO" → name the address: "...to elena.rostova@company.com (CEO)."
- Passive / artifact-centric phrasing - "The email mentions the storm" → "Agent mentions the storm in the email to Chloe."
- Tool-name annotations - drop "(via send_email)" and any tool names.
- Bundling independent actions into one rubric.
- Writing a Process rubric when a stricter Outcome would prove the same thing.
- Locking a Process rubric to one method/tool when the prompt named a goal.
- Overlapping rubrics that punish the same mistake more than once.
- Subjective language - "enough," "professional," "thorough," "helpful," "good."
- Using "approximately" for counts, IDs, or dates.
- Double-negative phrasing anywhere in a criterion.
- Rubrics the judge can't verify from the trajectory or final response.

11 Rubric quality checklist

Run this before you submit.

- ☐ The rubric exposes the model's failures - the agent does not pass every item.
 - ☐ Every criterion evaluates outcomes (with Process reserved for what Outcomes can't prove).
 - ☐ All critical events / required outputs have a rubric.
 - ☐ Every rubric has all three fields: criterion + justification + evidence.
 - ☐ Every criterion is agent-centric, self-contained, objective, atomic, and verifiable.
 - ☐ Repeated actions use spot checks (aggregate count + ≥ 3 specific instances), not one rubric per instance.
 - ☐ Multiple same-type write actions get one rubric per ground-truth item - not an "at least N" threshold.
 - ☐ Stacked rubrics use OR-logic correctly when multiple outcomes are valid.
 - ☐ Must-not-happen behaviors are penalty items (0/-1/-2/-3), written as positive statements; no double negatives.
 - ☐ Any must-not-happen behavior is covered by a penalty item.
 - ☐ Calculated/rounded numbers use "approximately"; counts, IDs, and dates use exact values.
 - ☐ Fuzzy values include examples + "(or similar)."
 - ☐ No rubric penalizes a valid alternative solution path.
 - ☐ Every NOT-PRESENT criterion is justified in all three areas - any that can't be are deleted.
 - ☐ Every important ask in the prompt is covered, with no big gaps or overlaps.
 - ☐ The gate (must-pass) set is small, objective, and covers the non-negotiables; nothing subjective is in the gate.
 - ☐ Every item declares category, grading_method, score type, weight, and dependencies.
 - ☐ Expert-Assessment items use distinct, falsifiable scored bands.
 - ☐ Penalties only subtract (0/-1/-2/-3); anything that must fail the run is a gate item instead.
-

Quick reference - key terms

Term	What it means
Objective Compliance	Rubric category — did the agent do the concrete, verifiable thing correctly (right action, content, facts)? Binary 0/1. Sub-types 1.1 / 1.2 / 2.1.
Expert Assessment	Rubric category — is the work actually good / would a senior professional trust it? Scored on the fixed -1...+1 band scale (this is where partial credit comes from).
Process	Rubric category — did the agent work soundly along the way, verified from the trajectory? Optional and rare.
Outcome	Collective shorthand for Objective Compliance + Expert Assessment (everything that isn't Process).
Gate (must-pass)	The small set of objective, non-negotiable criteria that decide Pass/Fail — all must pass, regardless of weight.
Partial credit	The weighted-average quality score from the non-gate criteria; diagnostic only, never decides Pass/Fail.
Band	A discrete score level on an Expert-Assessment item (from the fixed scale), each with a label + a falsifiable description.
Penalty	A subtract-only item (0 / -1 / -2 / -3) for must-not-happen behavior (e.g., fabrication)..
Weight	A positive integer (1–3) setting how much a criterion counts toward partial credit.
Per-item fields	The five fields every rubric declares: category, grading_method, score type, weight, dependencies.
grading_method	How a criterion is graded: deterministic (exact/code check) or LLM judge.
Score type	discrete / continuous / programmatic.

Dependencies	What the judge reads to grade an item: trajectory, produced artifact / final environment state, or a reference (gold source).
Pass@5	Difficulty metric — the share of runs (Gemini checkpoint, fully-specified prompt) over the model trajectories that clear the gate; must be under 30%.
Golden solution (GTFA)	The correct outcome (ground-truth final answer) you establish for the task (LLM assistance is fine); your reference for writing rubrics.
Agent trajectory	Collective term for the recorded agent runs (model, silver, and clean trajectories).
Model trajectory	The recorded agent run(s) on the fully specified prompt; where the Pass@5 difficulty gate is measured.
Silver trajectory	The recorded agent run on the under-specified prompt, guided with hints (expert-guided) — your source of truth for the rubric.
Clean trajectory	The silver trajectory reproduced without hints — the same trajectory run without the help of hints.